

# Final Project: Music Recommendation System

Music Informatics, Spring 2024

Informatics and Telecommunications, UoA

Konstantinos Chousos

Student ID: 1115202000215

Vasileios Katsaitis

Student ID: 1115202000073

Dimokritos Kolitsos

Student ID: 1115201900085

July 9, 2024

## Abstract

In this report, we present our implementation of a music recommendation system, based on a small dataset of 30 songs, consisting of three genre-related groups. We focus on the features related to their harmony, beat and timbre. From these features, we calculate the similarity of each song to the rest, in which we base our recommendations.

## Data Acquisition

Our dataset consists of 30 songs in total. Each one of us individually picked 10 songs for each of the genres ‘jazz’, ‘rock’ and ‘metal’. Every song is conventionally trimmed to the first 120 seconds (2 minutes), in order to reduce required dataset space and the model’s complexity. Also, each song is trimmed of any starting silence by removing any initial zero values in the audio array, so we get rid of useless information that might alter our extracted features negatively. Lastly, to illustrate our model’s results, every visualization will be provided for 3 selected songs, one from each genre — notated by index list `example_indices`.

## Feature Extraction

### Harmony and Melody

In this section, we examine each song in regards to its tonality and musical color. We focus on the scale used and its general pitch. We also examine separately the singer’s voice and the music generated solely from the instruments.

### Key detection

To find the key used in a song, we extract and examine the chromagram feature of it. To do that, we use the built-in function of `librosa` [1] `librosa.feature.chroma_stft`. To aggregate the values of each frame to create an average for each song, the mean was used, as it proved through experimentation to be more reliable.

To determine the key and mode (major or minor) of a song we used one of the key templates, such as Krumhansl’s [2] and Temperley’s [3], that were provided. Specifically, we found the best results to come from the usage of the `temperley05` template. Since the template was meant to correspond to the C major/minor scale, we use 10 other shifted variants of it to calculate the similarity to each key. This is done using the cosine similarity of the shifted template and the normalized chroma feature.

### Voice/source separation

For this task, the Meta’s Demucs library [4], [5] is used. We use the default pretrained “`htdemucs`” model and apply it to each song, separating between vocals and instrumental<sup>1</sup>. We also perform key detection in these newly acquired instrumental audio signals to examine any differences between them and the original songs. Some differences were observed, but not enough to warrant further examination.

---

<sup>1</sup>To be more accurate, the model actually separates each audio file to ‘vocals’, ‘drums’, ‘bass’ and ‘other’. Such segmentation was decided to be redundant and thus the instruments were combined to result in a general ‘instrumental’ part.

## Pitch detection

To estimate a fundamental frequency/pitch for each of the songs, the YIN algorithm [6] is used — specifically its librosa implementation<sup>2</sup>. For each song, we acquire a time series of fundamental frequencies in Hertz. To calculate the representative fundamental frequency for the whole song, we find the most common one.

The pitch of the standalone instrumentals and vocals is also calculated. The difference between the original’s pitch and the instrumental’s for each song is negligible, but the vocal’s fundamental frequency is, as expected, quite higher.

## Rhythm and Tempo

### Onset Detection

Implementation-wise, we use the **spectral-based** method to calculate the Novelty Curve, as it covers a wider range of songs (which are not particularly percussive, for example) and, in general, better detects changes in the spectral content of a song.

In order to calculate the novelty curve, we need to compute the magnitude spectrum of input signal, which we do using **STFT**. For sample rate = 44100 Hz we set the corresponding **STFT parameters** as follows: Number of samples (N) = 2048 and hop length = 512. This choice resulted from multiple reruns of our algorithm, as these yielded the more accurate tempo predictions (corresponding to those of the librosa tempo estimation function). The way we verified the validity of the parameter values was by setting different values to the prementioned parameters (e.g. N = 4096, 1024, ... and hop length = 1024, 256, ...) and calculating the averaged tempo differences between them. We picked that combination of parameters, which yields the smallest tempo deviations. Furthermore, the final parameter initializations we chose are compatible with the parameter initialization in Muller’s code (where for sample rate = 22050 Hz he sets N = 1024 and hop length = 256).

In the **logarithmic compression** stage of the algorithm, we experimented with different hyperparameter C (gamma) values [1, 10, 100]. It was decided that the optimal value of the hyperparameter, which ensures a sufficient level of compression while preserving adequate information about energy change

in the spectrum, is C = 100. We, therefore, did not need to experiment with more values.

The novelty curve is, then, calculated after applying all the necessary steps of the spectral-based method (compression - normalization) and by also subtracting, from it, its **local average**. When calculating the local average we heavily rely on Muller’s formula and code [7], which uses a moving average filter. This post-processing technique leads to a smoother and more precise novelty curve for each audio file.

### Tempo Extraction

To finally extract the optimal tempo value from the novelty curve, we use the **autocorrelation** technique on it. This was a conscious choice, as the majority of songs in our dataset consist of jazz and rock songs, in which a listener typically measures the tempo in measure or tactus level (half-time).

After applying autocorrelation to the novelty curve, we conventionally set a **range of acceptable bpm values**. Thus, in the event that a tempo value does not belong to that range, we are forced to look for other harmonics of it that belong to these permissible limits. This range will be rather large, e.g. from 30 to 350 bpm.

### Beat Tracking

The final tempo value returned by our algorithm (**manual approach**) results as that tempo value, within the allowed limits, for which the maximum autocorrelation value is found. This value is then **compared** with the corresponding tempo value returned by the **librosa function**, where it’s decided, depending on their distance, which tempo value is finally returned as the most valid. By default, we consider librosa function results as the optimal ones, however if tempo deviation is small enough, we return the value of the manual approach. This distance (deviation) is defined as 1 bpm.

As far as the parameters are concerned, we did not need to experiment with the acceptable tempo range, as almost every result fell into that range. Specifically, we only set the **start\_bpm** parameter as 65 (librosa function), since from the range of values we experimented with [60, 65, 70, 75, 80], this was the minimum value possible with which we had the most valid tempo predictions. However, for songs

<sup>2</sup><https://librosa.org/doc/latest/generated/librosa.yin.html>

belonging to the 'metal' genre, **metal\_start\_bpm** = 100 is used as the starting bpm value. This design choice was made, because typically in a metal song a half-time tempo is not intuitive enough, as is perhaps the case in a jazz or rock song. So, we get a double-time tempo value which is considered more appropriate for a metal song.

### Plotting

We conventionally plot the **novelty curve**, **tempogram**, **autocorrelogram** and song **waveform** with pulses as the 4 most important graphs related to task 3B. We assume that these graphs are sufficient enough to better understand how our approach works, since our implementation does not rely on any commonly used audio features, such as MFCC's, pitch and spectral features (which are actively used in the other tasks).

### Timbre & Spectral Shape

For starters I didn't really have a grasp of what Timbre represents and after going back to the lecture PDFs and also doing some extensive research on the internet I must confess that I still don't fully get it.. Nevertheless I have a pretty good idea now at least.

#### MFCCs

I begun by extracting MFCCs from each song. I didn't have to do anything fancy since there is already a librosa function which extracts MFCCs from audio. I started with 13 coefficients as the task description dictates and is the most common value if you want to get the most relevant information to the spectral envelope. I then experimented with taking the mean, median and standard deviation for each song to see what results it yields and mess around with them because I was still not sure of what would be the best feature and in what form to keep it for the recommendation system to use as a parameter. We decided on 3 songs, 1 from each genre that we would do all our plots with so I then plotted the MFCCs heatmap for each song and then a graph of their mean, median and std deviation values for each coefficient, to visualize the differences between both the features and the songs from each genre, which would be helpful in understanding if a certain feature would be more helpful to use for the recommendation system. I then tried out the same

process but with a different number of coefficients each time but I realized that the only difference was in the information added from the extra coefficients (obviously..) so I decided to just plot the above but with 20 coeffs just to get the whole picture. Since plotting the data of just 3 song didn't sum up the whole picture and also later they were not as useful in a rough implementation of a recommendation function I wrote using only the spectral features, I wanted to check the differences of each genre regarding the mean, median and standard deviation to see if it was even worth bothering using them as features. I did that by grouping by and finding the mean of each genre for each feature. The results were pretty much the same for mean and median and a little different for std deviation but the I came to the outcome that the were not any significant differences after all.. maybe that has to do with the genres we selected and that some of the songs we selected in particular do not have such pronounced differences. Also at least for the mean and median I came to the conclusion that after 14 coeffs there was not much to separate the 3 genres so for the final feature to be used regarding mfccs I kept the mfccs up until 14 coeffs and also added the std deviation that I had for 20 coeffs (which we ended up not using anyway). Later to try and make the MFCCs feature more pronounced and representative I added to it the 1st and 2nd derivatives to make up for a more comprehensive feature that would yield better results in the recommendation system.

#### Timbre descriptors (spectral centroid, spectral bandwidth, and spectral rolloff)

For other timbre descriptors after reading the lecture slides and doing my own research I decided on extracting spectral centroid (brightness), spectral bandwidth (noise/tonal attribute) and spectral rolloff(brightness/ spectral distribution). After extracting them I decided to plot for each feature in a single plot all 3 of the selected songs to have a clear visualization of the differences a song from each genre has regarding each feature. I then added the features to the final dataframe which also contains all the above mentioned features that I saved in it to use in the recommendation system.

Then to do some tests and figure out if the features I kept are helpful and how helpful they are for a recommendation system I plotted a similarity matrix using

cosine similarity (I also implemented a recommend songs function to run test and this helped me decide on which features were not really helping and then to either re-evaluate how I extract them or not use them at all and most importantly it helped in deciding which features were actually usefull). Seeing the similarity matrix we can see clearly that at least the metal songs have very similar timbral attributes, the jazz songs also kinda make for a group or 2 with similarities and finally the rock songs are not that well defined at least from the perspective of timbre only features. But we can get better results if we combine the extracted features from Harmony/Melody and Rhythm & Tempo as well.

## Recommendation System

Our system for the recommendation of similar tracks is based on a *similarity matrix* computed from the resulting feature matrix of the dataset — for this task, we used the cosine similarity metric as the most suitable one. In our code, we leveraged the Sklearn library [8]. Through trial and error, we found that the combination of the

1. song's key/scale,
2. song's pitch/fundamental frequency,
3. song's tempo,
4. the mean of the MFCCs and
5. the mean of other spectral features

was the one that led to the best results. We also experimented with using the instrumental's and vocals' key and pitch, but these features hadn't much of a positive impact. Also, the mode of the songs is discarded, as it is a categorical feature and thus didn't differentiate them enough.

Our recommendation system provides the  $K$  most similar tracks and the  $K$  most *dissimilar*. A usage example is the following:

```
import random

query_index = random.randint(0, len(dataset) - 1)
recommend_songs(query_index)
```

And the corresponding output:

```
3 nearest recommendations for the song: Gojira
↳ - Adoration for none
```


Nearest Recommendations:

Fit For An Autopsy - Black Mammoth  
Lamb of God - Blood of the Scribe  
Lamb of God - Again We Rise

Furthest Recommendations:

Snarky Puppy - Shofukan (We Like It Here)  
Ryo Fukui - Scenery  
Miles Davis - So What

A visualization of our dataset is shown in fig. 1.



**Figure 1:** The 3D plot of our dataset. The 3 dimensions are the song's key, its tempo and the mean MFCC extracted from it.

Another approach that was tested was the usage of Spotify's Voyager library [9] for the creation of the similarity space and the  $K$ -NN graph. This approach was limited in regards to the retrieval of the most dissimilar tracks and thus was discarded. Also the K-Means algorithm was shortly examined as an alternative, but it presented the same issue.

## Conclusion

In this report, we presented our music recommendation system implementation, that focuses on melodic, rhythmic and spectral features for each similarity

comparisons. Our results vary between the genres, in large part due to common characteristics. For example, metal tracks are most distinguishable from the others, due to their fast beat and spectral maximalism. In contrast, some jazz songs were found to be similar with some rock ones. At first glance this would seem to be a flaw of our system, but by listening to the query song and the recommendations we came to the conclusion that they *do* actually share a lot of elements and can be thought of as neighbors in our small dataset.

## Task separation

- Harmony and Melody feature exploration was done by Konstantinos Chousos.
- Rhythm and Tempo feature exploration was done by Vasileios Katsaitis.
- Timbre and Spectral Shape feature exploration was done by Dimokritos Kolitsos.

## References

- [1] B. McFee, M. McVicar, D. Faronbi, *et al.*, *Librosa/librosa: 0.10.2.post1*, version 0.10.2.post1, Zenodo, May 14, 2024. DOI: [10.5281/ZENODO.11192913](https://doi.org/10.5281/ZENODO.11192913). [Online]. Available: <https://zenodo.org/doi/10.5281/zenodo.11192913>.
- [2] C. L. Krumhansl, *Cognitive Foundations of Musical Pitch* (Cognitive Foundations of Musical Pitch). New York, NY, US: Oxford University Press, 1990, pp. x, 307, x, 307, ISBN: 978-0-19-505475-0 978-0-19-514836-7.
- [3] D. Temperley, “What’s Key for Key? The Krumhansl-Schmuckler Key-Finding Algorithm Reconsidered,” *Music Perception*, vol. 17, no. 1, pp. 65–100, Oct. 1, 1999, ISSN: 0730-7829. DOI: [10.2307/40285812](https://doi.org/10.2307/40285812). [Online]. Available: <https://online.ucpress.edu/mp/article/17/1/65/62051/What-s-Key-for-Key-The-Krumhansl-Schmuckler-Key>.
- [4] *Facebookresearch/demucs*, Meta Research, Jul. 8, 2024. [Online]. Available: <https://github.com/facebookresearch/demucs>.
- [5] S. Rouard, F. Massa, and A. Défossez, “Hybrid transformers for music source separation,” in *Icassp* 23, 2023.
- [6] A. De Cheveigné and H. Kawahara, “YIN, a fundamental frequency estimator for speech and music,” *The Journal of the Acoustical Society of America*, vol. 111, no. 4, pp. 1917–1930, Apr. 1, 2002, ISSN: 0001-4966, 1520-8524. DOI: [10.1121/1.1458024](https://doi.org/10.1121/1.1458024). [Online]. Available: <https://pubs.aip.org/jasa/article/111/4/1917/547221/YIN-a-fundamental-frequency-estimator-for-speech>.
- [7] M. Müller and A. Villar-Corrales. “Spectral-Based Novelty.” (), [Online]. Available: [https://www.audiolabs-erlangen.de/resources/MIR/FMP/C6/C6S1\\_NoveltySpectral.html#Spectral-Based-Novelty](https://www.audiolabs-erlangen.de/resources/MIR/FMP/C6/C6S1_NoveltySpectral.html#Spectral-Based-Novelty).
- [8] F. Pedregosa, G. Varoquaux, A. Gramfort, *et al.*, “Scikit-learn: Machine learning in Python,” *Journal of Machine Learning Research*, vol. 12, pp. 2825–2830, 2011.
- [9] *Spotify/voyager*, Spotify, Jul. 7, 2024. [Online]. Available: <https://github.com/spotify/voyager>.